

Saumya Parag Phadkar

saumya.phadkar@gmail.com [Portfolio](#) [LinkedIn](#) [GitHub](#) New York City

Experience

Machine Learning Engineer

Sept 2025 – Present

Sanio AI, California, USA

- Built a LangChain + OpenRouter-based Dynamic Attribute Extraction Pipeline enabling adaptive routing.
- Implemented hybrid normalization with embeddings + fuzzy logic with 90% precision.
- Designed a schema-free EAV framework with attribute coverage scoring for catalog analytics.
- Deployed FastAPI microservices orchestrated via Temporal + Docker on AWS ECS.

Software Engineer

Sept 2025 – Present

Streetlives, Inc., NY, USA

- Built hybrid multilingual chat flow combining LLM slot-filling and guided options for text + voice interaction.
- Replaced free-form LLM replies with a SQL-driven dialogue engine for accurate service recommendations.
- Integrated RPA-based form automation and privacy-safe CCaaS routing to streamline social-service intake.

Senior Engineer and Research Analyst

Oct 2024 – Sept 2025

New York University, NY, USA

- Designed NYU's first centralized PhD Alumni Career Outcomes Dashboard on Tableau Cloud.
- Benchmarked 50 R1 universities' dashboards to design NYU's metric framework for outcome reporting.
- Led accessibility compliance achieving WCAG 2.1 AA standards and improved usability.

Education

New York University, New York, NY

Sept 2023 – May 2025

MS in Computer and Information Sciences

Savitribai Phule Pune University (University of Pune), Pune, India

Aug 2019 – Jun 2023

BE in Computer Engineering, Hons in Artificial Intelligence and Machine Learning

Projects

SustainableShopper: AI-Powered Virtual Wardrobe & Shopper

[GitHub Repo](#) 

JWT, MongoDB, .NET, CLIP, User Research, UI/UX,

May 2025

- Built Flask backend with JWT auth, MongoDB, and async polling to Fashn.ai for virtual try-on inference.
- Integrated CI/CD pipelines using Docker & Jenkins, sustaining 98% uptime across 200+ real-time sessions.
- Developed React dashboards & user flows, connecting .NET Core APIs for e-commerce interoperability.
- Deployed GPT-4 + CLIP multimodal outfit recommendation system; achieved 91% user preference rate.

Robust Stock Movement Prediction from Noisy Small Data

[GitHub Repo](#) 

Python, PyTorch, Meta-Learning, LSTM, Siamese Networks, Learn2Learn

May 2024

- Simulated noise, sparsity & feature loss to evaluate model robustness on Apple, Microsoft, etc. stock data.
- Applied MAML meta-learning & Siamese networks, improving F1-stability under degraded data.
- Benchmarked LSTM, KNN, SVM & neural models using Optuna cross-validation & hyperparameter tuning.

UrbanHealthOpt: Optimization of Urban Health Facility Placement

[GitHub Repo](#) 

Big Data, Python, PyTorch, scikit-learn, TensorFlow, Scipy

Dec 2024

- Built a geospatial simulator with TIGER shapefiles to optimize placement of 300+ urban hospitals.
- Integrated demographic and epidemiological datasets from U.S. Census and ACS.
- Designed custom loss functions to run 50+ simulations, presenting data-driven insights for urban planning.

Technical Skills

Domains: AI/ML/DL, GenAI, RAG, Prompt Engineering, Explainable AI, Data Viz, DevOps, Cloud

Programming Languages: Python, SQL, JavaScript, Java, C++, C, R, Shell

Databases: MySQL/MSSQL, PostgreSQL, MongoDB, Spark, Neo4j, OntoText, Hadoop, Oracle

Technologies: Git, AWS, Azure, Docker, Temporal, .NET Core, Tableau, Power BI, FastAPI

Frontend: HTML, CSS, JavaScript, TypeScript, React, Redux, Django, Angular, Vue3, Flask, D3

AWS Certified AI Practitioner: AWS Services, AI Ethics, AI Governance, Secure AI, Trustworthy AI